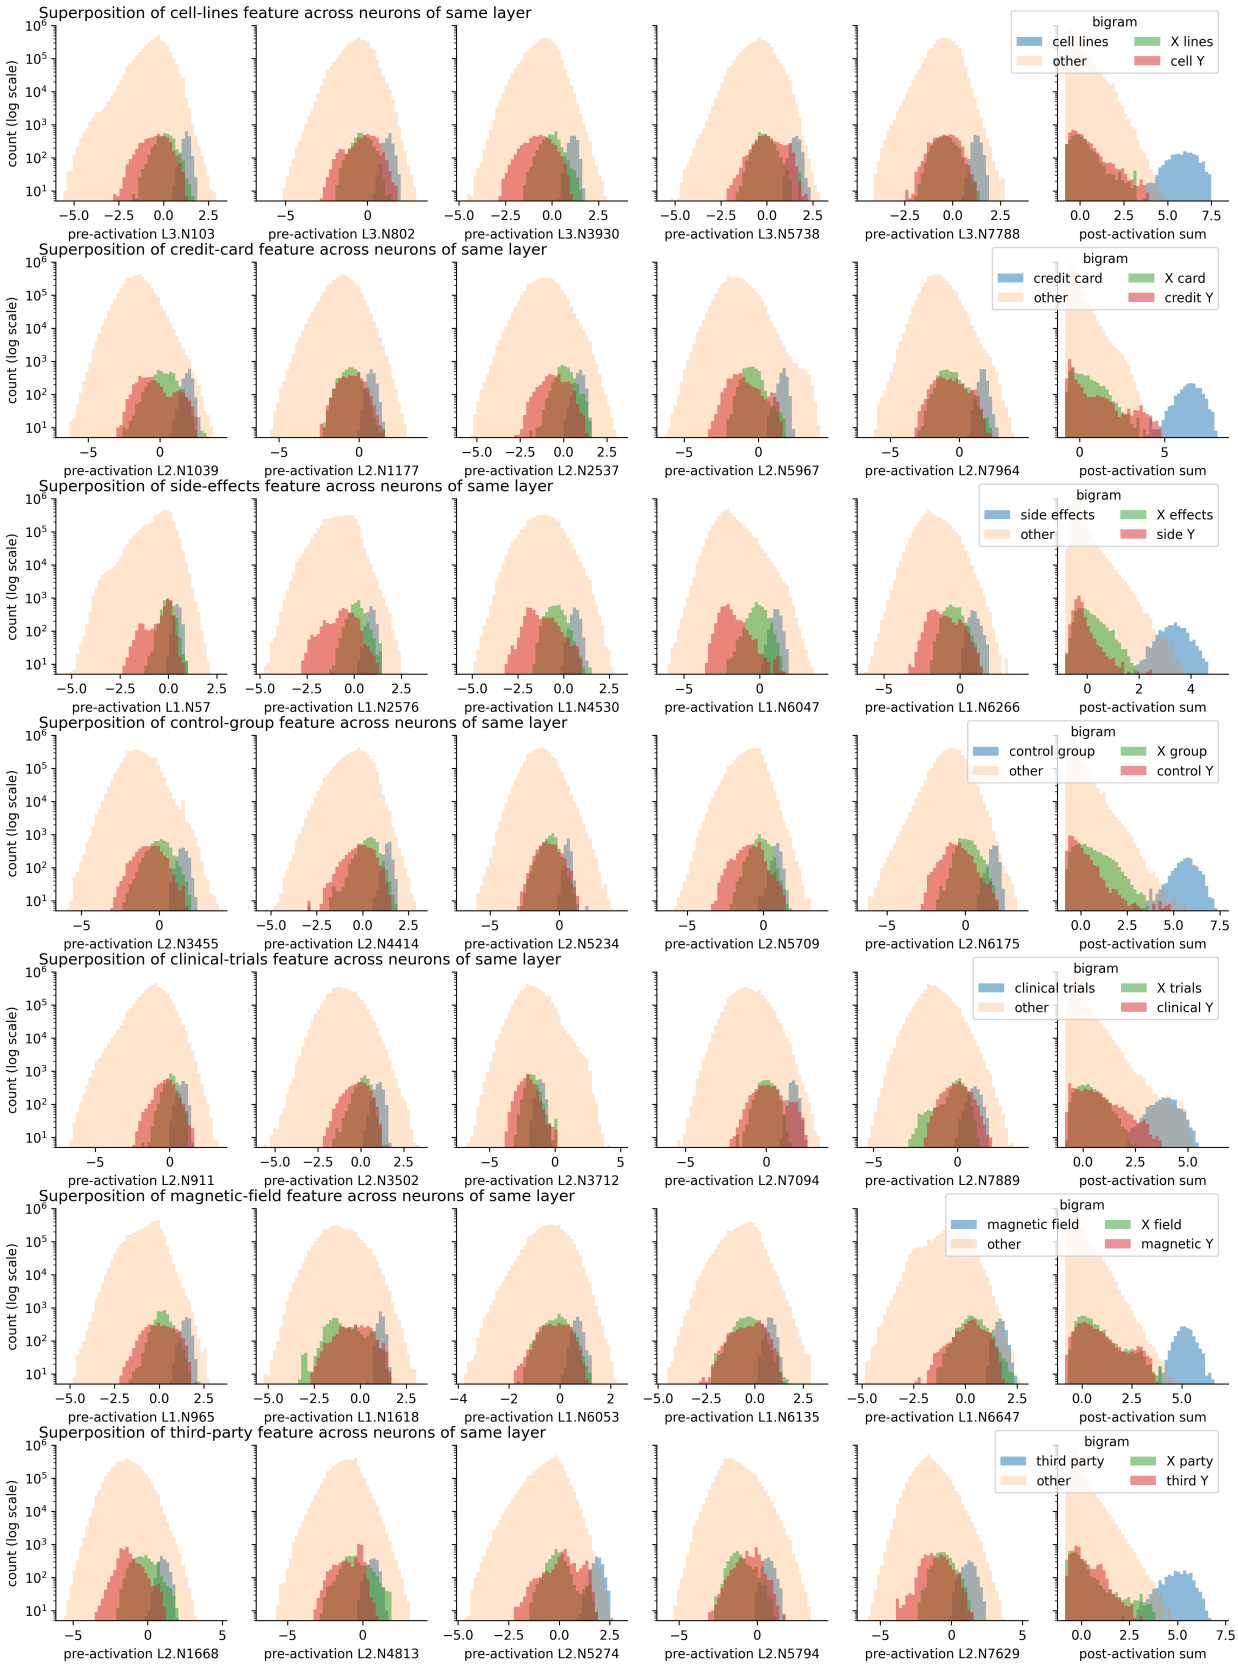

Figure 11: More examples of superposition of compound words in Pythia-1B

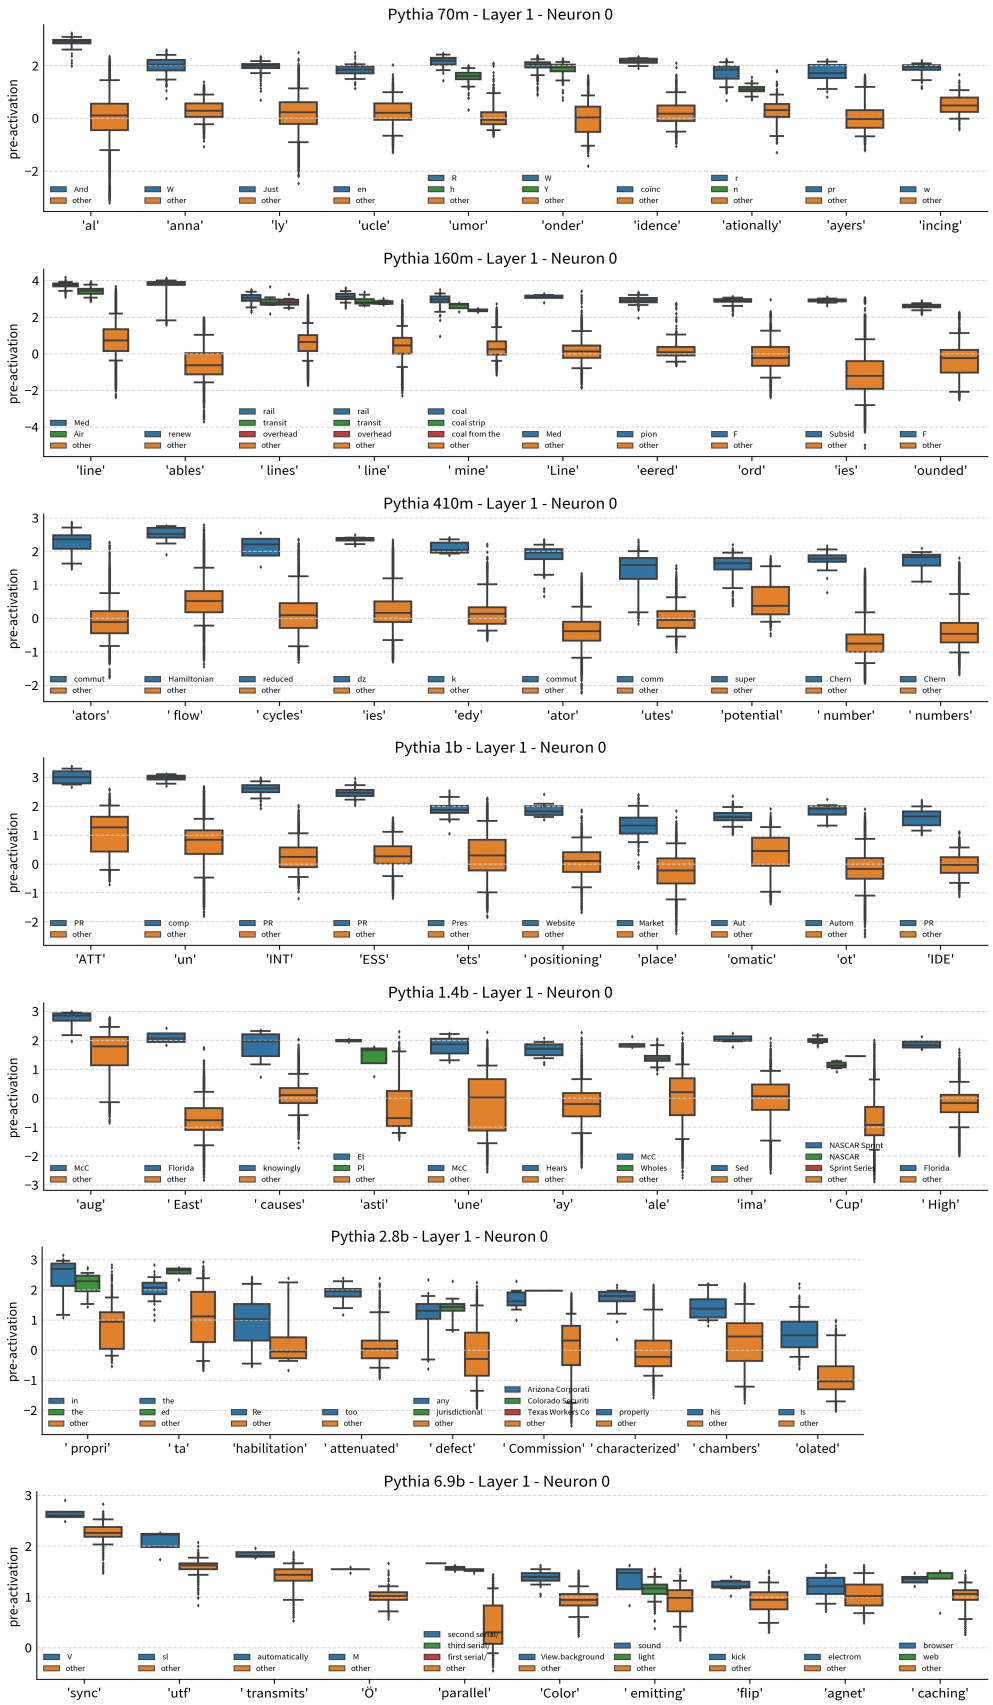

Figure 12: More examples of n-gram polysemanticity in layer 1 neuron 0 of all Pythia models studied.

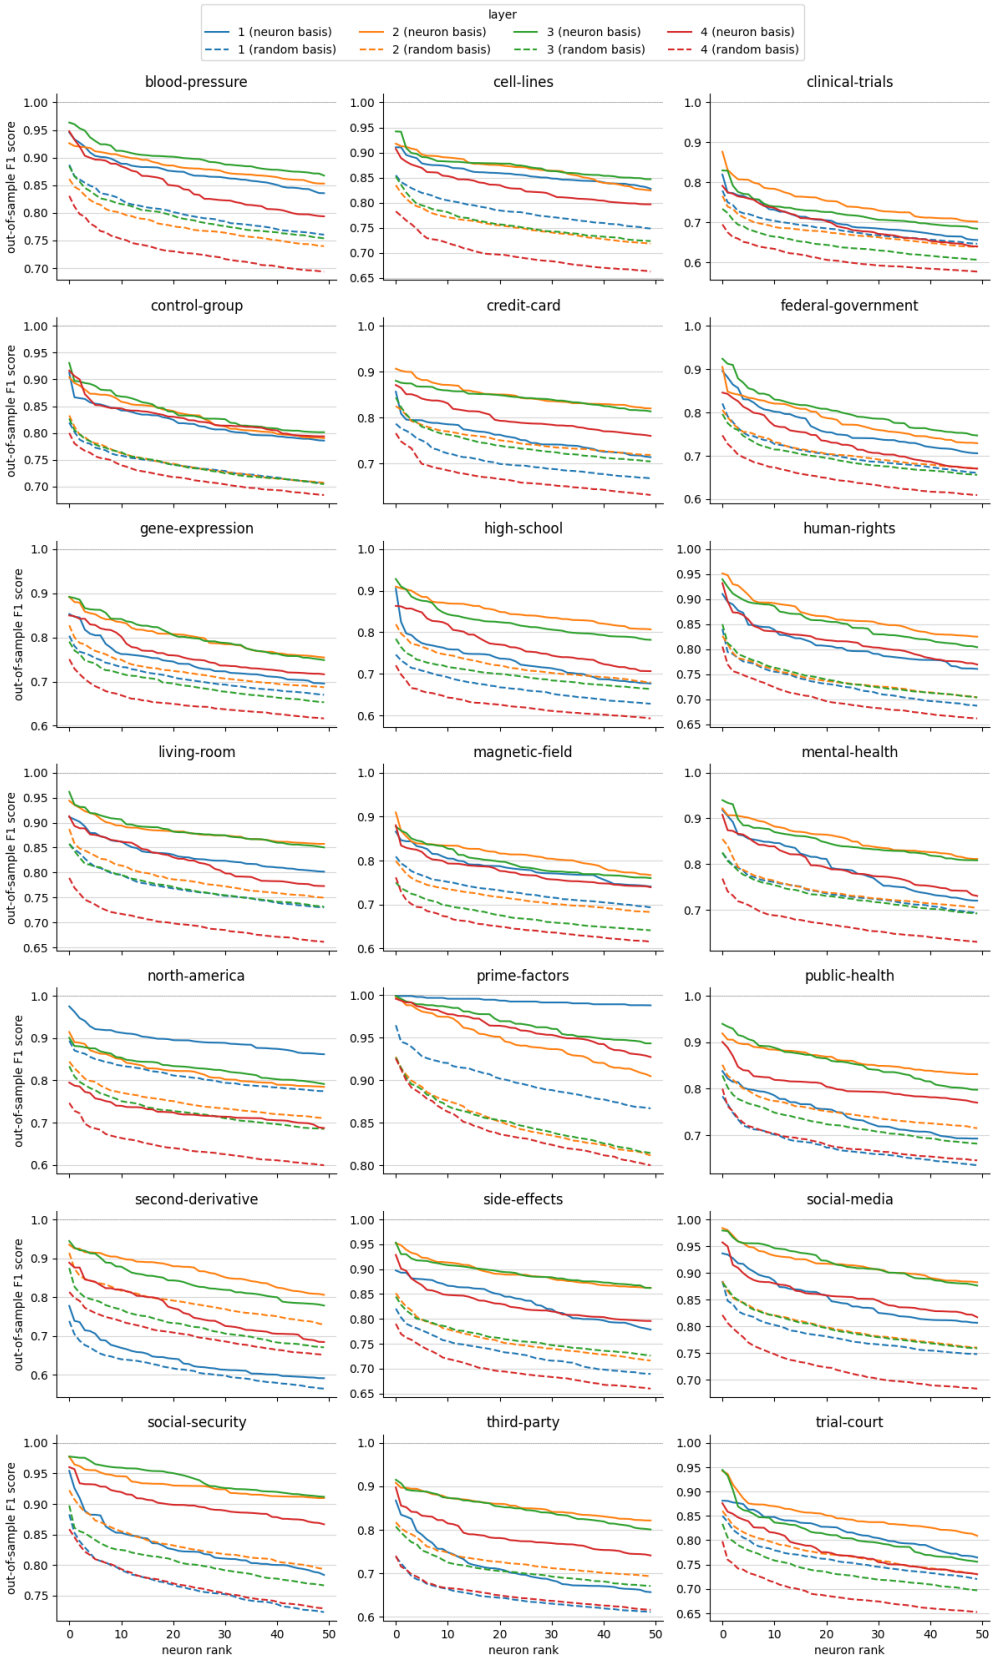

Figure 13: Demonstration of neuron basis alignment in early layers of Pythia-1B. Compares the compound-word classification performance of the top 50 neurons for layers 1-4 in the standard neuron basis, as opposed to a random basis (the post-activation dataset multiplied by a random Gaussian $d \times d$ matrix).